



INDIGO - DataCloud

RIA-653549



INDIGO-Datacloud

WP4.1 - Update on Computing Virtualization

Alvaro Lopez Garcia ¹

¹aloga@ifca.unican.es

Spanish National Research Council - CSIC



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



Summarized (DoW) Roadmap

- **MS19 (Month 15):** Support for container execution with trusted container repository and standard interfaces.
- **MS22 (Month 15):** Implementation of spot-instances mechanisms
- **MS20 (Month 23):** Execution of containers through batch systems.
- **MS21 (Month 30):** Proof of concept for accessing hardware accelerators (GPUs) and low latency networks (IB) in containers.
- **MS23 (Month 30):** Implementation of Advanced Scheduling Policies
- **MS24 (Month 30):** Delivery of the orchestration engine with standard interfaces whenever possible (TOSCA)

Where are we?

- Month 8th.
- Detailed roadmap and current progress tracked in OpenProject:
<https://project.indigo-datacloud.eu/projects/wp4/timelines/16>



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



General remarks on container support I

- Docker is the container technology chosen.
- Objective: Execute containers as light-weight virtual machines.
- Try to not introduce API changes, treat them as normal resources from the Cloud Management Framework (CMF) point of view.
- Try to offer **similar functionality** on both CMFs.
- Higher level features are outside WP: orchestration, application management done in other WPs.



General remarks on container support II

- OpenStack initial considerations:
 - LXC support in-tree, officially supported.
 - OpenStack nova-docker ¹ driver available out-tree.
 - OpenStack Magnum: new component for managing containers (Kubernetes).
 - ▶ Gaining a lot of momentum.
 - ▶ Overlaps with WP5 components.
 - ▶ We will follow up the developments, but not focusing on them.

- OpenNebula initial considerations:
 - Driver available only for LXC.
 - No Docker driver.
 - Several approaches to dcker integration presented in ONE Conf. in Barcelona ².
 - Megam project, similar to OpenStack Magnum.

¹<https://github.com/openstack/nova-docker>

²<http://opennebula.org/docker-machine-and-opennebula/>, slides: <http://goo.gl/2kA9uv>



Container Support in OpenStack

- Evaluation floating document <https://goo.gl/SViKGj>.
 - Evaluation of LXC driver, following up development.
 - Evaluation of nova-docker driver, following up development.
 - ▶ Some functionality missing: e.g. attach/detach volume.
 - Evaluation of Magnum project will start next month.
- Openstack + Docker testbeds deployed at LIP and INFN-PD, currently under testing
- No API changes introduced (nova-docker, LXC), Magnum implies an API change.



Container Support in OpenNebula

Summary

13

1. Docker as a hypervisor ⇒ NO
2. Distribute OpenNebula in Docker ⇒ NO
3. **Integrate with Docker-Machine ⇒ WIP**
4. **OneFlow + Docker ⇒ Medium-term Roadmap**

- ONE is focusing on Docker, but using a different approach (see slide on the left).
- No Docker hypervisor driver foreseen.
- **WP4.1 is developing it, filling one of the gaps in ONE development.**
- Initial implementaiton of **onedock** driver available
<https://github.com/indigo-dc/onedock>.
- No API changes introduced.

© OpenNebula Systems SL

OpenNebula.systems

Slide from OpenNebula Conference in
Barcelona, October 2015
<http://goo.gl/2kA9uv>



Docker-related issues

- Not possible to attach a volume (i.e. block device) in runtime.
- Do we need this? Not possible in Docker, therefore not possible in both ONE and OpenStack.
- Functionality is implemented, developers are doing bikeshedding: they do not agree on the CLI.
- If we need this we need to push for this functionality.
- OpenStack (nova-docker) and Docker related bugs:
 - <https://bugs.launchpad.net/nova-docker/+bug/1321817>
 - <https://github.com/docker/docker/pull/8348>
 - <https://github.com/docker/docker/pull/8826>
 - <https://github.com/docker/docker/issues/8829>



Other related issues

- OpenStack nova only makes possible to run one driver in each compute node.
- Therefore, running nova-docker needs a dedicated compute node, excluding those nodes from running other VMs.
- Initial work on running two hypervisor drivers in OpenStack (e.g. KVM/Xen + nova-docker) in one host.
- If the functionality is required by several resource providers we should formalize this and submit it to OpenStack.



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



Repository integration

- Work depends on ONE and OpenStack Docker driver implementations.
- Activity will start next month.



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



OCCI container support

- Starting testing of OCCI + OpenStack + nova-docker.
- Small changes foreseen in current implementation.
- Available testbed by end of year (end of M9).



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



- Synergy being already tested and INFN-PD.
- Unfunded contribution expected from CENSNET student.
- Waiting for the instructions on how to proceed with is repo (OpenStack CI – INDIGO sync).
- Initial problems with the SLURM priority multi-factor algorithm.
 - Scheduling algorithm being re-adapted.
 - This work will take about 1 month.
- ONE integration with existing prototype not possible due to ONE architecture.
 - Requirements collected³ and new design available⁴.
 - Implementation phase started.

³https://project.indigo-datacloud.eu/work_packages/685

⁴https://project.indigo-datacloud.eu/work_packages/610



Spot-instances support

- Blueprint submitted to OpenStack developers⁵.
 - Feedback received, initial design slightly modified.
 - Deciding how to expose functionality.
 - ▶ Via API extension, introducing API changes.
 - ▶ Via flavors, no API changes.
 - ▶ Via tags, no API changes.
- Implementation has started regardless of BP status, taking into account OpenStack devs suggestions.
- Changes in the implementation foreseen, **not in the functionality**.
- Code will be published as a fork inside INDIGO namespace before end of November.
- Deciding how to expose the spot-instances functionality.
- Initial prototype being tested at CSIC testbed.

⁵<https://review.openstack.org/#/c/104883/>



Partition director

- Implementation started.
- Work stalled a bit, as Stefano had a car accident.



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



Batch system integration

- Two different approaches being tested.
 - epilog/prolog -based
 - nodocker/shipper being more promising.
- Docker with CUDA 7.5 and pyOpenCL.
- Testing together job submission + GPGPUS together.
- Planning to use the CAFFEE CUDA neural network.
- Writing the procedure for running those applications.

WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



TOSCA in OpenStack

- Orchestration is done through OpenStack HEAT.
- TOSCA parser + TOSCA to HOT translator based support.
- The translator is a CLI tool supporting only HOT translation so far.
- Possible issues doing the translation at the API level.
- We had a meeting with IBM people (the team responsible of the TOSCA tools).
 - Planning to have regular meetings with them.
 - IBM people is developing a translation service, details are coming in the following weeks.



TOSCA in OpenNebula

- Work has been shifted 6 months back, already started, aligned with OpenStack task.
- Orchestration support is done through the Infrastructure Manager (IM).
- TOSCA parser + TOSCA to RADL translator based support.
- RADL target language plugin will be developed as a plugin for the existing translator.
 - The same tool will be used for both translations.
- Implementation already started



- Parser is being extended to support INDIGO requirements.
- Initial work started by UPV.
- Fork already available at <https://github.com/indigo-dc/tosca-parser>.
- Need to check how INDIGO changes can be pushed/integrated to upstream parser.
 - General changes must be pushed upstream.
 - What about INDIGO-only related changes?
 - Need to check how we can extend the parser for those cases where integration is not possible.



WP4.1 - Container support

WP4.2 - Repository integration

WP4.3 - OCCl extension to support containers

WP4.4 - Improved scheduling

WP4.5 - Integration of cont. in batch systems

WP4.6 - IaaS TOSCA orchestration

Conclusions



General remarks

- Developers demand for clear procedures on how to comply with WP3 requirements.
- Code needs to be pushed in a controlled way into <http://github.com/indigo-dc>.
 - Specially if it is going to be pushed to upstream projects.
 - Ensure things merged are compliant with INDIGO QA (WP3.1).
- We will be releasing pilots from next month.
 - Communication channels with WP3 for deployment on pilot testbeds.
 - Communication channels with WP4, WP5 for testing.



Next steps and ETAs

- M9: OpenStack + nova-docker + OCCl pilot.
- M9: ONE + onedock + OCCl pilot.
- Madrid 2nd architecture meeting
- M10: Spot instances pilot prototype.
- M10: Fair-share scheduling pilot prototype.
- M11: Automatic sync with INDIGO DockerHub at the site-level.
- Q1 2016: initial TOSCA orchestration service available.

Indigo Web: <http://indigo-datacloud.eu>
aloga@ifca.unican.es